

# Introduction to Shapefiles

## Introduction

One of my more recent areas of growth as a data scientist has been learning to work with spatial data. If you have worked with yield monitor data, *ArcGIS*, or used an application map to vary the application rate of your fertilizer, you have worked with spatial data. We will learn how to plot spatial data in detail in Unit 12, but here is a brief introduction to importing and inspecting spatial data.

## R Packages

To look at spatial data, we will use the R package **sf** (short for *spatial features*). Remember, when we use a new package in R for the first time, we must install and launch it with the following commands. The **sf** package takes a while to install, so it is already installed in *RStudio Cloud*. As a reminder, however, the installation code is included below. In this case it has a hashtag (#) in front of it. The hashtag converts code to a **comment**; any code that follows a hashtag on the same line will not be run. This is a useful way to block code we don't want to run, or add an explanation about code we are running

Go ahead and restart R and clear the output, as you learned in the *summary statistics* exercise, then install and load the **sf** package with the code below:

```
# install.packages('sf')
library(sf) # this loads the sf package


## Registered S3 methods overwritten by 'tibble':
##   method      from
##   format.tbl pillar
##   print.tbl  pillar

## Linking to GEOS 3.8.0, GDAL 3.0.4, PROJ 6.3.1; sf_use_s2() is TRUE
```

## Reading Shapefiles

To read in a shape file, we use the **st\_read()** function. We read in a shape file the same way we read in a .csv file in the previous exercise. We follow **st\_read()** with the path to the shape file, which has a .shp suffix.

If you view a shape file in the file directory, it is usually accompanied by identically named files with different extensions (the three letters after the last period). For example, if you navigate to the directory *Cloud>project>exercises>data>Agriculture\_Census-shp* in the window in the lower right, you will see the *Agriculture\_Census.shp* file, but also *Agriculture\_Census.prj* and three other files. These additional files tell R how to draw the shapefile as a map. Reading in the shape file (*Agriculture\_Census.shp*) will cause these other files to be read in automatically.



We will read in a shape file from Indiana that includes county summaries for a couple of dozen agricultural statistics.

```
indiana = st_read("data/Agriculture_Census-shp/Agriculture_Census.shp")

## Reading layer `Agriculture_Census` from data source
##   `/cloud/project/exercises/data/Agriculture_Census-shp/Agriculture_Census.shp'
##   using driver `ESRI Shapefile'
## Simple feature collection with 92 features and 31 fields
## Geometry type: POLYGON
## Dimension:     XY
## Bounding box: xmin: -88.08975 ymin: 37.77306 xmax: -84.78768 ymax: 41.76233
## Geodetic CRS:  WGS 84
```

## Examining Spatial Feature Data Frames

Let's inspect the first six rows of the spatial feature data frame using the `head()` command, where the name of the data frame is given between the parentheses.

```
head(indiana)

## Simple feature collection with 6 features and 31 fields
## Geometry type: POLYGON
## Dimension:     XY
## Bounding box: xmin: -87.52704 ymin: 41.16195 xmax: -84.80631 ymax: 41.76233
## Geodetic CRS:  WGS 84
##   OBJECTID          AGCENSPO20          COUNTY N_FARM TOTCROP_AC HAVCROP_AC DOL_FARM
## 1      1            2515    Elkhart County    1335     160160     143324     92912
## 2      2            2516 Saint Joseph County    666      139661     131004     82849
## 3      3            2517    La Porte County    749      226816     212940     127922
```

```

## 4      4    2518    Steuben County    581    99218    71944    44133
## 5      5    2519    Lagrange County   1392   156233   127875   74194
## 6      6    2539    Lake County     442    138929   132551   108206
##   NURS_DOL LVSTOK_DOL SALE_2500 SALE_5000 EXPEN_1000 NURSY_1000 N_CHICKENS
## 1    34688    89350    251     100    96862     514   3005354
## 2    40051    15127    144      71    40479     583     852
## 3    68041    27773    136      60    71666     742   -9999
## 4    16776    8865     236      44    20288      0   -9999
## 5    37364    65915    223     105    83668      0   1870836
## 6    42747    5080     101      30    34712     66      0
##   N_CATTLE N_PIGS CORN_AC WHEAT_AC COTTON_AC SOYBEAN_AC VEGTABL_AC ORCHARD_AC
## 1    42719   73951   64955    3861      0   42251    1338      0
## 2    6440    27430   69251    4073      0   45696     480      0
## 3    24980   27110  113242    4186      0   76809    1125      0
## 4    9257    6859    32152    4483      0   25120     304      0
## 5    39275   69338   61262    3993      0   32666    1182      0
## 6    3204    9435   68344    3101      0   55698    1256      0
##   PECAN_AC IRRGTD_AC N_MACHINE FEMALE_OP MINORTY_OP FAM_FARM igs_DBO_Ag
## 1      0    23524     495     32      0   1164      0
## 2      0    12941     444     44      3   584       0
## 3      0    27090     631     44      0   626       0
## 4      0    1225    -9999     42      0   505       0
## 5      0    23478     315     26      0   1232      0
## 6      0    6211     327     42      3   356       0
##   SHAPE_Leng SHAPE_Area           geometry
## 1  139058.4 1210349270 POLYGON ((-86.06246 41.7619...
## 2  155098.8 1196846667 POLYGON ((-86.22555 41.7614...
## 3  175216.0 1563031845 POLYGON ((-86.82619 41.7608...
## 4  116305.4 835438144 POLYGON ((-85.19779 41.7605...
## 5  128976.5 1000327995 POLYGON ((-85.29134 41.7606...
## 6  167477.6 1300778609 POLYGON ((-87.52401 41.7135...

```

Spatial feature data frames are very similar to the data frames that were introduced to in the last exercise. There are a few differences, however. First, every spatial feature data frame consists of rows of observations that are linked, in the geometry column, with a georeference. That georeference may be a simple point on the map, or it may be a polygon that defines an area, such as a field, a county, or a state.

Spatial feature data frames also include some metadata about the geometry of the file, including geometry type (POLYGON in this case, since we are working with county outlines) and – this is critical – *CRS*. CRS stands for **coordinate reference system**, the “projection” we use to draw a map, and the georeferences used to define a place or area.

Maps are flat representations of a round object, and there are different ways to draw these maps. We need not go into a deep review of those projects at this point; suffice it to know that each projection represents tradeoffs in how accurately it portrays areas in the center and periphery of the map and how much it represents the curvature. For most of my work, I use the CRS of 4326. You know this system more commonly as the latitude/longitude system.

The bounding box defines where the map is centered. Latitude/longitude maps are more accurate in portraying areas near the center of the map than the periphery, so it is important the map be centered on the area of interest.


Don’t worry about committing the above information to memory for now, we will discuss projection is much greater detail in Unit 12.

## Visualizing Data

To visualize the data, we can use the `plot()` function from the `sf` package. You might assume we could map a column of data the same way we visualized its contents earlier, by indicating the data frame and column name, separated by the dollar sign. We will try that with `TOTCROP_AC`, the total number of acres per county in crop production.

```
tot_crop = indiana$TOTCROP_AC
```

```
plot(tot_crop)
```



Instead of a map, however, this shows us the range of values for `TOTCROP_AC` in the dataset.

To view a map, we need to first subset the data a little differently than we did before: we can't use the dollar-sign approach. That only assigns the values of our single column to our new R object.

See what happens if we look at the top of our `tot_crop` object using the `head()` command

```
head(tot_crop)
```

```
## [1] 160160 139661 226816 99218 156233 138929
```

The object only has a single column of numbers with the county values. It lacks information on the geometry of those counties, so that they can be drawn.

To draw our map, we need to assign two columns to our new object: `TOTCROP_AC` and geometry columns. We do this by referencing the desired measure as `indiana["TOTCROP_AC"]`. This tells R to pass both the `TOTCROP_AC` values and their geometries to the new object, `tot_crop`.

```
tot_crop = indiana["TOTCROP_AC"]
head(tot_crop)
```

```
## Simple feature collection with 6 features and 1 field
## Geometry type: POLYGON
## Dimension: XY
```

```

## Bounding box: xmin: -87.52704 ymin: 41.16195 xmax: -84.80631 ymax: 41.76233
## Geodetic CRS: WGS 84
##   TOTCROP_AC           geometry
## 1    160160 POLYGON ((-86.06246 41.7619...
## 2    139661 POLYGON ((-86.22555 41.7614...
## 3    226816 POLYGON ((-86.82619 41.7608...
## 4    99218  POLYGON ((-85.19779 41.7605...
## 5   156233  POLYGON ((-85.29134 41.7606...
## 6   138929  POLYGON ((-87.52401 41.7135...


```

When we examine the top six rows of our object, tot\_crop, we see it contains both the county values and the county geometries.

Note: we only use this approach, with the desired data column in quotes and brackets, with shape file data. The rest of the time, the data frame and desired column should be separated by the dollar sign, as we learned above.

We can then map these data:

```
plot(tot_crop)
```



No one can resist a pretty map. The technical name for a map where political shapes like counties are color-coded, is *chloropleth*. We saw plenty of these in 2020, between the pandemic and the election year.


Back to happier thoughts. Let's look at the mean farm income per county. This is indicated by the "DOL\_FARM" column. Again, we first subset the data and then plot it.

```

farm_income = indiana[["DOL_FARM"]]
plot(farm_income)

```


## DOL\_FARM



It looks like the larger farms, by income, are in northwest Indiana, just south of Lake County and the Chicago suburbs. Where are the chickens in Indiana? Let's find out:

```
chickens = indiana["N_CHICKENS"]
plot(chickens)
```

## N\_CHICKENS




It would appear the most finger-lickin' good county is Elkhart County, in the far north. Think of it as a big "cluck you" to Michigan, which is immediately north of those counties. Appropriately, there are also a couple of counties near Kentucky!

## Create Your Own Maps

Where are the hogs in Indiana? Create a map for "N\_PIGS".

```
chickens = indiana["N_CHICKENS"]
plot(chickens)
```

## N\_CHICKENS



Where are most of the wheat acres? Create a map for “WHEAT\_AC”.

Which counties have greater farm ownership by minorities? Create a map for “MINORTY\_OP” (note the missing “I”).